

# Machine Learning

## Practice 2

Quan Minh Phan & Ngoc Hoang Luong

University of Information Technology (UIT)

April 7, 2021

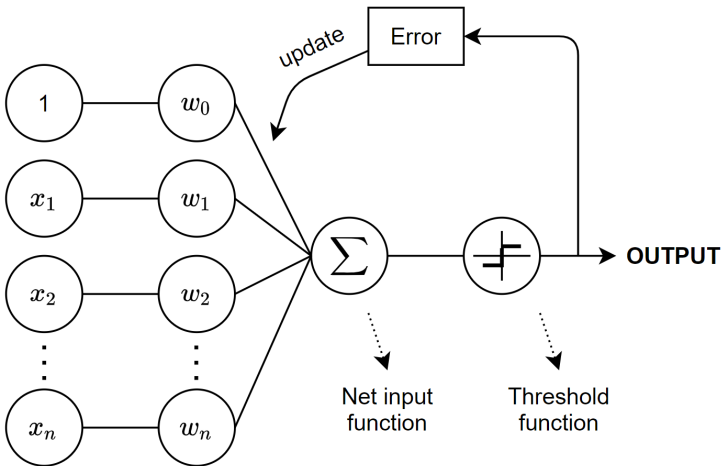
# Table of contents

1 Perceptron

2 Linear Regression

3 **Adaptive Linear Neuron (Adaline)**

4 Logistic Regression



**Figure:** The general concept of the perceptron

## Problem

Using the perceptron model to classify the species of flowers ('setosa' or 'versicolor') based on the sepal and petal width.

- 2 features: the sepal width  $x_1$ ; the petal width  $x_2$   
 $x = [x_1, x_2]$
- 2 labels: 'setosa'  $\leftarrow -1$ ; 'versicolor'  $\leftarrow 1$
- $w = [w_1, w_2]$
- Net input function ( $z$ )  
 $z = w_1 * x_1 + w_2 * x_2$
- Unit step function ( $\phi(\cdot)$ )


$$\phi(z) = \begin{cases} 1 & \text{if } z \geq \theta, \\ -1 & \text{otherwise.} \end{cases}$$

# Recall (next)


- $w = [w_0, w_1, w_2]$
- Net input function ( $z$ )  
 $z = w_0 + w_1 * x_1 + w_2 * x_2 = w^T x$
- Unit step function ( $\phi(\cdot)$ )

$$\phi(z) = \begin{cases} 1 & \text{if } z \geq 0, \\ -1 & \text{otherwise.} \end{cases}$$

# Unit Step Function



$$f(x) = \begin{cases} 1 & x \geq 0, \\ -1 & \text{otherwise.} \end{cases}$$



$$f(x) = \begin{cases} 0.5 & x \geq 1, \\ -1 & \text{otherwise.} \end{cases}$$

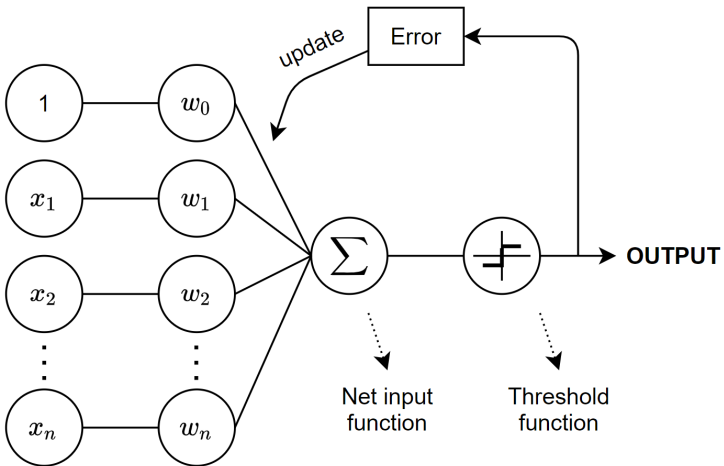


Figure: The general concept of the perceptron

# Training process

---

**Algorithm 1** Pseudocode for the training process

---

```
1: Initialize the weights,  $w$ 
2: while Stopping Criteria is not satisfied do
3:   for  $x \in X$  do
4:     Compute the output value,  $\hat{y}$ 
5:     Updates the weights
6:   end for
7: end while
```

---



# Updating the weights

- $w = w + \Delta w$
- $\Delta w_i = \eta * (y - \hat{y}) * x_i$   
where:
  - ▶  $\eta$ : learning rate
  - ▶  $y$ : the true class label
  - ▶  $\hat{y}$ : the predicted class label

## Examples

$$\Delta w_0 = \eta * (y - \hat{y})$$

$$\Delta w_1 = \eta * (y - \hat{y}) * x_1$$

$$\Delta w_2 = \eta * (y - \hat{y}) * x_2$$

# Components

## Hyperparameters

- $\eta \rightarrow$  the learning rate
- $\text{max\_iter} \rightarrow$  the maximum number of epochs
- $\text{random\_state} \rightarrow$  to make the reproducible results

## Parameters

- $w \rightarrow$  the weights of model
- $\text{errors} \rightarrow$  to store the error in each epoch

## Methods

- $\text{fit}(X, y) \rightarrow$  to train the model
- $\text{predict}(X) \rightarrow$  to predict the output value
- $\text{net\_input}(X) \rightarrow$  to combine the features with the weights

# Implement (code from scratch)

```
class Peceptron_:
```

```
    def __init__(self, eta = 0.01, max_iter = 20, random_state = 1):  
        self.eta = eta  
        self.max_iter = max_iter  
        self.random_state = random_state  
        self.w = None  
        self.errors = [ ]
```

```
    def net_input(self, X):  
        return np.dot(X, self.w[1:]) + self.w[0]
```

```
    def predict(self, X):  
        return np.where(self.net_input(X) ≥ 0.0, 1, -1)
```

```

def fit(self, X, y):
    rgen = np.random.RandomState(self.random_state)
    self.w = rgen.normal(loc = 0.0, scale = 0.01, size = 1 + X.shape[1])
    self.errors = [ ]
    for n_iter in range(self.max_iter):
        n_wronglabels = 0
        idx = rgen.permutation(len(y))
        X, y = X[idx], y[idx]
        for xi, yi in zip(X, y):
            error = yi - self.predict(xi)
            self.w[1:] += self.eta * error * xi
            self.w[0] += self.eta * error
            n_wronglabels += int(error != 0.0)
        self.errors.append(n_wronglabels)

```

# Implement (library)

```
from sklearn.linear_model import Perceptron
```

## Hyperparameters

- eta
- max\_iter
- random\_state

## Parameters

- coef\_
- intercept\_

## Methods

- fit( $X, y$ )
- predict( $X$ )

# Practice

- Using 'iris.csv' dataset
- How can we use the 'sepal length' and 'sepal width' to classify the speices of flower?


# Data visualization

```
>> import matplotlib.pyplot as plt
```

```
>> idx_setosa = y_train == -1  
    idx_versicolor = y_train != -1
```

```
>> plt.scatter(X_train[idx_setosa, 0], X_train[idx_setosa, 1], color='red',  
               marker='s', label='setosa')  
    plt.scatter(X_train[idx_versicolor, 0], X_train[idx_versicolor, 1],  
               color='blue', marker='x', label='versicolor')  
    plt.xlabel('sepal length [cm]')  
    plt.ylabel('sepal width [cm]')  
    plt.legend(loc='upper left')  
    plt.show()
```

# Data visualization





# Practice

```
>> ppn = Perceptron_(eta = 0.001, max_iter = 30, random_state = 1)  
ppn.fit(X_train, y_train)
```


```
>> from sklearn.linear_model import Perceptron
```

```
>> ppn_ = Perceptron(eta = 0.001, max_iter = 30, random_state = 1)  
ppn_.fit(X_train, y_train)
```

# Plotting the errors

```
>> plt.plot(range(1, len(ppn.errors) + 1), ppn.errors, marker='o')  
plt.xlabel('Epochs')  
plt.ylabel('# Missifications')  
plt.show()
```

# Plotting the errors



# Practice

```
>> w_ppn = ppn.w  
w_ppn
```


```
>> [-0.01575655  0.06508244 -0.11108172]
```

```
>> w_ppn_ = np.append(ppn_.intercept_, ppn_.coef_)  
w_ppn_
```

```
>> [-0.006  0.0196 -0.0325]
```

```
>> plot_decision_regions(X_train, y_train, classifier=ppn)
plt.xlabel('sepal length [cm]')
plt.ylabel('sepal width [cm]')
plt.legend(loc='upper left')
plt.show()
```


# Visualization



# Practice (next)


- Create a new model and train it on data after standardizing

# Data visualization

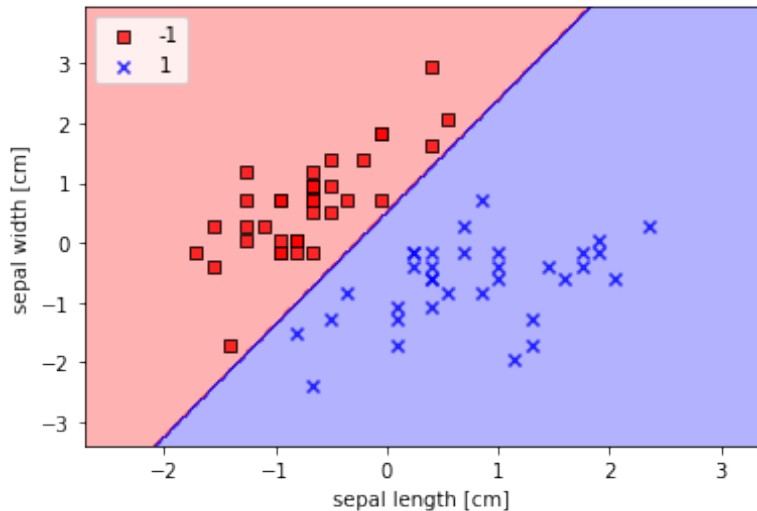




# Plotting the costs



# Plotting the results



# Table of contents

1 Perceptron

2 Linear Regression

3 Adaptive Linear Neuron (Adaline)

4 Logistic Regression

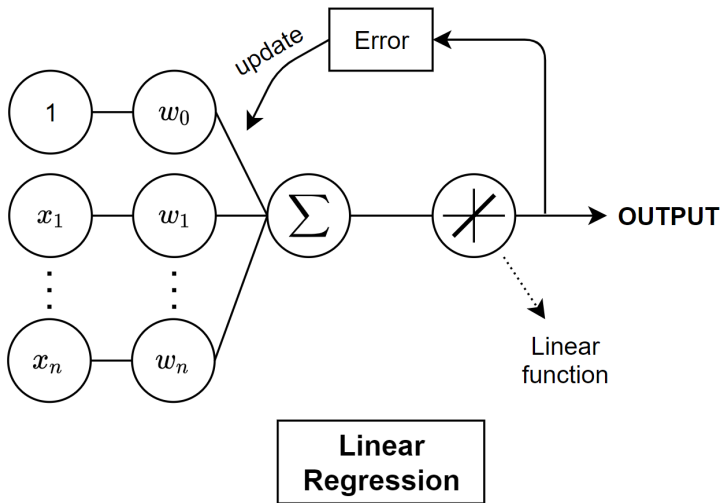


Figure: The general concept of Linear Regression

# Minimizing cost functions with gradient descent

Cost function:

$$J(w) = \frac{1}{2} \sum_i (y^{(i)} - \phi(z^{(i)}))^2$$

Update the weights:

$$w := w + \Delta w$$

$$\Delta w = -\eta \nabla J(w)$$

$$\frac{\partial J}{\partial w_j} = - \sum_i (y^{(i)} - \phi(z^{(i)})) x_j^{(i)}$$

$$\Delta w_j = -\eta \frac{\partial J}{\partial w_j} = \eta \sum_i (y^{(i)} - \phi(z^{(i)})) x_j^{(i)}$$

# Minimizing cost functions with gradient descent

$$w_j = \begin{cases} w_j + \eta * X^T \cdot \text{dot}(y - \phi(z)) & j \in [1, \dots, n] \\ w_j + \eta * \text{sum}(y - \phi(z)) & j = 0 \end{cases}$$

# Pseudocode of Training process

---

## Algorithm 2 Gradient Descent

---

- 1: Initialize the weights,  $w$
  - 2: **while** Stopping Criteria is not satisfied **do**
  - 3:   Compute the output value,  $\hat{y}$
  - 4:   Updates the weights
  - 5: **end while**
-

# Components

## Hyperparameters

- eta
- max\_iter
- random\_state

## Parameters

- w
- costs

## Methods

- fit( $X, y$ )
- predict( $X$ )
- net\_input( $X$ )



# Implement (code from scratch)

```
class LinearRegression_GD:
```

```
    def __init__(self, eta = 0.001, max_iter = 20, random_state = 1):  
        self.eta = eta  
        self.max_iter = max_iter  
        self.random_state = random_state  
        self.w = None  
        self.costs = [ ]
```

```
    def net_input(self, X):  
        return np.dot(X, self.w[1:]) + self.w[0]
```

```
    def predict(self, X):  
        return self.net_input(X)
```

```

def fit(self, X, y):
    rgen = np.random.RandomState(self.random_state)
    self.w = rgen.normal(loc = 0.0, scale = 0.01, size = 1 + X.shape[1])
    self.costs = [ ]
    for n_iters in range (self.max_iter):
        errors = y - self.predict(X)
        self.w[1:] += self.eta * X.T.dot(errors)
        self.w[0] += self.eta * errors.sum()
        cost = (errors**2).sum() / 2
        self.costs.append(cost)

```

# Implement (library)

## Stochastic Gradient Descent

```
from sklearn.linear_model import SGDRegressor
```

### Hyperparameters

- eta0
- max\_iter
- random\_state

### Parameters

- intercept\_
- coef\_

### Methods

- fit(X, y)
- predict(X)

# Implement (library)

## Normal Equation

```
from sklearn.linear_model import LinearRegression
```

### Parameters

- intercept\_
- coef\_

### Methods

- fit(X, y)
- predict(X)

# Differences

## Gradient Descent

- $w := w + \Delta w$
- $\Delta w = \eta \sum_i (y^{(i)} - \phi(z^{(i)})) x^i$

## Stochastic Gradient Descent

- $w := w + \Delta w$
- $\Delta w = \eta (y^{(i)} - \phi(z^{(i)})) x^i$

## Normal Equation

- $w = (X^T X)^{-1} X^T y$


# Practice

- Using 'housing.csv' dataset
- How can we use the 'average number of rooms' (RM) to estimate the 'price' of houses (MEDV)?

# Plotting data

```
>> plt.scatter(X_train, y_train, c='steelblue', edgecolor='white', s=70)
plt.xlabel('Average number of rooms [RM] (standardized)')
plt.ylabel('Price in $1000s [MEDV] (standardized)')
plt.show()
```

# Practice





# Practice

## Gradient Descent

```
>> reg_GD = LinearRegression_GD(eta=0.001, max_iter=20,  
    random_state=1)  
reg_GD.fit(X_train, y_train)
```

## Stochastic Gradient Descent

```
>> reg_SGD = SGDRegressor(eta0=0.001, max_iter=20,  
    random_state=1, l1_ratio=0, tol=None, learning_rate='constant')  
reg_SGD.fit(X_train, y_train)
```


## Normal Equation

```
>> reg_NE = LinearRegression()  
reg_NE.fit(X_train, y_train)
```

# Plotting the cost

```
>> plt.plot(range(1, len(reg_GD.costs) + 1), reg_GD.costs)
plt.xlabel('Epochs')
plt.ylabel('Cost')
plt.title('Gradient Descent')
plt.show()
```

## Gradient Descent



# Practice

```
>> w_GD = reg_GD.w  
w_GD
```

```
>> [0.00767139  0.64623542]
```

```
>> w_SGD = np.append(reg_SGD.intercept_, reg_SGD.coef_)  
w_SGD
```

```
>> [0.00783841  0.64551218]
```

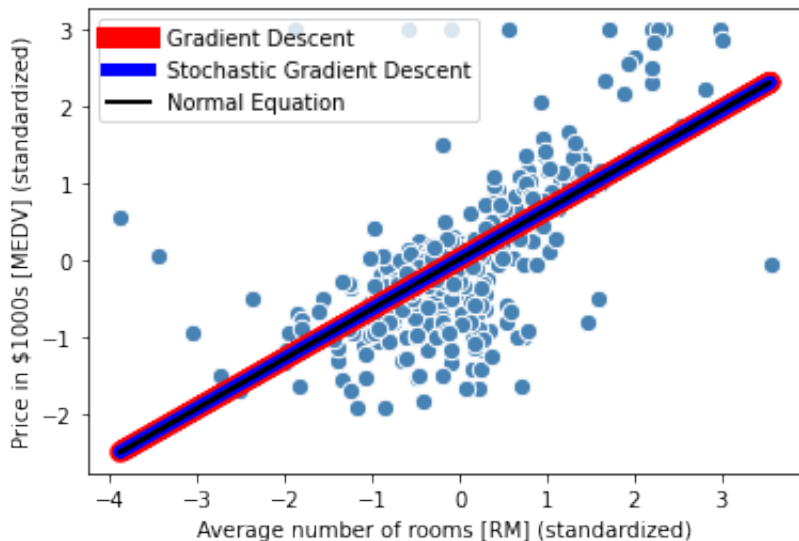
```
>> w_NE = np.append(reg_NE.intercept_, reg_NE.coef_)  
w_NE
```

```
>> [0.00773059  0.64638912]
```

# Plotting the results

```
>> plt.scatter(X_train, y_train, c='steelblue', edgecolor='white', s=70)
plt.plot(X_train, reg_GD.predict(X_train), color='red', lw=10,
label='Gradient Descent')
plt.plot(X_train, reg_SGD.predict(X_train), color='blue', lw=6,
label='Stochastic Gradient Descent')
plt.plot(X_train, reg_NE.predict(X_train), color='black', lw=2,
label='Normal Equation')
plt.xlabel('Average number of rooms [RM] (standardized)')
plt.ylabel('Price in $1000s [MEDV] (standardized)')
plt.legend()
plt.show()
```

# Plotting the results



# Practice

```
>> y_pred_1 = reg_GD.predict(X_test)
```

```
>> y_pred_2 = reg_SGD.predict(X_test)
```

```
>> y_pred_3 = reg_NE.predict(X_test)
```

# Performance Evaluation

```
>> from sklearn.metrics import mean_absolute_error as MAE  
    from sklearn.metrics import mean_squared_error as MSE  
    from sklearn.metrics import r2_score as R2
```

## Mean Absolute Error

```
>> print('MAE of GD:', round(MAE(y_test, y_pred_1), 6))  
    print('MAE of SGD:', round(MAE(y_test, y_pred_2), 6))  
    print('MAE of NE:', round(MAE(y_test, y_pred_3), 6))
```

## Mean Squared Error


```
>> print('MSE of GD:', round(MSE(y_test, y_pred_1), 6))  
    print('MSE of SGD:', round(MSE(y_test, y_pred_2), 6))  
    print('MSE of NE:', round(MSE(y_test, y_pred_3), 6))
```



## $R^2$ score

```
>> print('R2 of GD:', round(R2(y_test, y_pred_1), 6))  
      print('R2 of SGD:', round(R2(y_test, y_pred_2), 6))  
      print('R2 of NE:', round(R2(y_test, y_pred_3), 6))
```

# Learning rate too large



# Polynomial Regression


## Example

```
X = [258.0, 270.0, 294.0, 320.0, 342.0, 368.0, 396.0, 446.0, 480.0, 586.0]  
y = [236.4, 234.4, 252.8, 298.6, 314.2, 342.2, 360.8, 368.0, 391.2, 390.8]
```

```
>>> X = np.array([258.0, 270.0, 294.0, 320.0, 342.0, 368.0, 396.0, 446.0,  
480.0, 586.0])[ :, np.newaxis]  
y = np.array([236.4, 234.4, 252.8, 298.6, 314.2, 342.2, 360.8, 368.0,  
391.2, 390.8])
```

```
>>> plt.scatter(X, y, label='Training points')  
plt.xlabel('X')  
plt.ylabel('y')  
plt.legend()  
plt.show()
```

# Plotting data



# Polynomial Regression

```
>> from sklearn.linear_model import LinearRegression  
lr = LinearRegression()  
lr.fit(X, y)
```

# Polynomial Regression

## Syntax

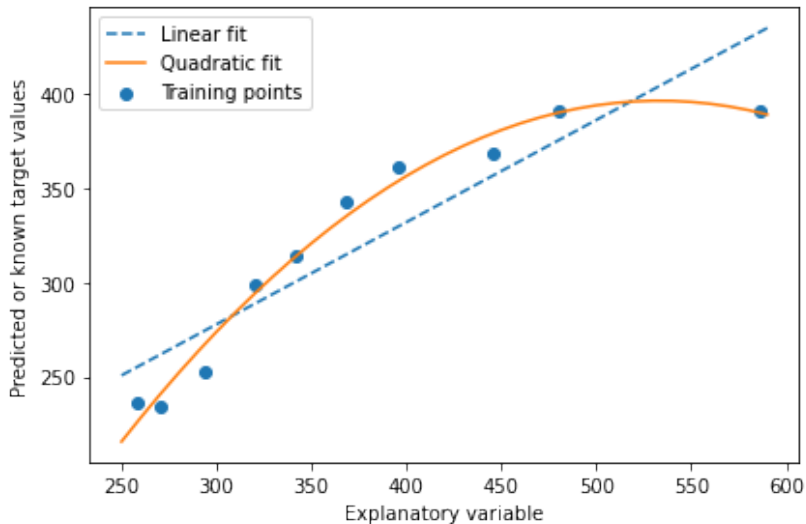
```
from sklearn.preprocessing import PolynomialFeatures
```

```
>> from sklearn.preprocessing import PolynomialFeatures  
pr = LinearRegression()  
quadratic = PolynomialFeatures(degree=2)  
X_quad = quadratic.fit_transform(X)  
pr.fit(X_quad, y)
```

```
>> X_fit = np.arange(250, 600, 10)[: , np.newaxis]
```

```
>> y_fit_linear = lr.predict(X_fit)  
y_fit_quad = pr.predict(quadratic.fit_transform(X_fit))
```

```
>> plt.scatter(X, y, label='Training points')
plt.xlabel('X')
plt.ylabel('y')
plt.plot(X_fit, y_fit_linear, label='Linear fit', linestyle='--')
plt.plot(X_fit, y_fit_quad, label='Quadratic fit')
plt.legend()
plt.tight_layout()
plt.show()
```





# Practice

## Linear regression

```
>> lr = LR()  
    lr.fit(X_train, y_train)
```

## Polynomial regression (quadratic)

```
>> quadratic = PolynomialFeatures(degree=2)  
    X_quad = quadratic.fit_transform(X_train)  
    pr_quad = LR() pr_quad = pr_quad.fit(X_quad, y_train)
```

## Polynomial regression (cubic)

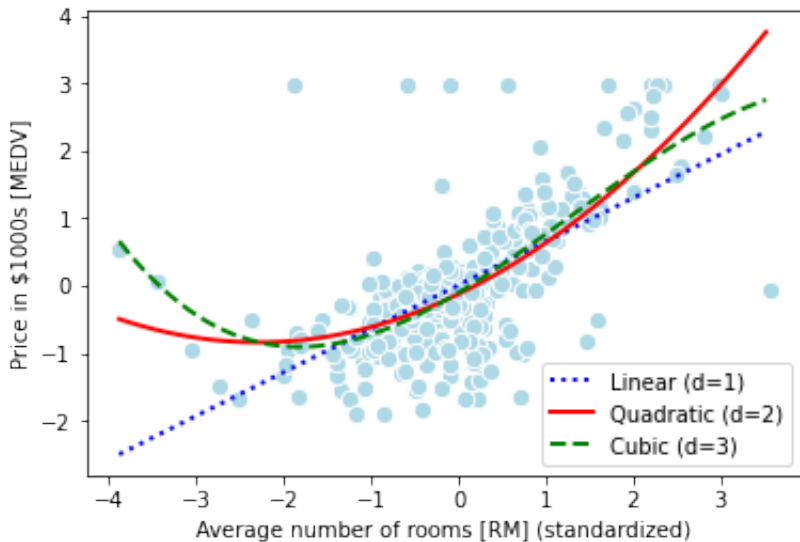
```
>> cubic = PolynomialFeatures(degree=3)  
    X_cubic = cubic.fit_transform(X_train)  
    pr_cubic = LR() pr_cubic = pr_cubic.fit(X_cubic, y_train)
```

```
>> X_fit = np.arange(X_train.min(), X_train.max(), 0.1)[: , np.newaxis]
```

```
>> y_linear_fit = lr.predict(X_fit)  
y_quad_fit = pr_quad.predict(quadratic.fit_transform(X_fit))  
y_cubic_fit = pr_cubic.predict(cubic.fit_transform(X_fit))
```

## Plotting the results

```
>> plt.scatter(X_train, y_train, c='steelblue', edgecolor='white', s=70)
plt.plot(X_fit, y_lin_fit, label='Linear (d=1)', color='blue', lw=2,
linestyle=':')
plt.plot(X_fit, y_quad_fit, label='Quadratic (d=2)', color='red', lw=2,
linestyle='-')
plt.plot(X_fit, y_cubic_fit, label='Cubic (d=3)', color='green',
lw=2, linestyle='-')
plt.xlabel('Average number of rooms [RM] (standardized)')
plt.ylabel('Price in $1000s [MEDV] (standardized)')
plt.legend()
plt.show()
```



# Table of contents

- 1 Perceptron
- 2 Linear Regression
- 3 Adaptive Linear Neuron (Adaline)**
- 4 Logistic Regression

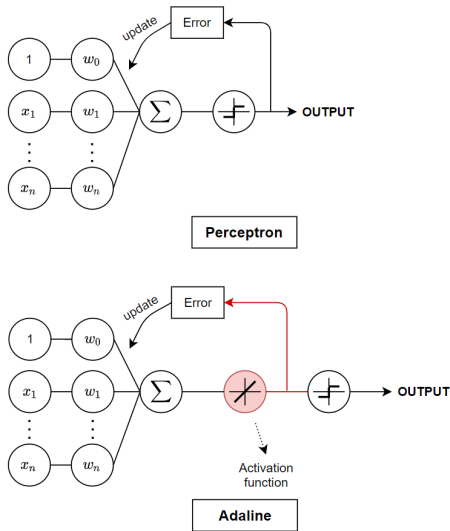


Figure: Differences between Perceptron and Adaline

# Training process

---

**Algorithm 3** Pseudocode for the training process

---

```
1: Initialize the weights,  $w$ 
2: while stopping criteria is not satisfied do
3:   for  $x \in X$  do
4:     Compute the output value,  $\hat{y}$ 
5:     Updates the weights
6:   end for
7: end while
```

---

# Updating the weights

- $w = w + \Delta w$
- $\Delta w_i = \eta * (y - \hat{y}) * x_i$   
where:
  - ▶  $\eta$ : learning rate
  - ▶  $y$ : the true class label
  - ▶  $\hat{y}$ : the predicted class label

## Examples

$$\Delta w_0 = \eta * (y - \hat{y})$$

$$\Delta w_1 = \eta * (y - \hat{y}) * x_1$$

$$\Delta w_2 = \eta * (y - \hat{y}) * x_2$$



# Components

## Hyperparameters

- eta
- max\_iter
- random\_state

## Parameters

- $w$
- costs

## Methods

- $\text{fit}(X, y)$
- $\text{predict}(X)$
- $\text{net\_input}(X)$
- $\text{activation}(X)$

# Implement (code from scratch)

```
class Adaline:
```

```
    def __init__(self, eta = 0.01, max_iter = 50, random_state = 1):  
        self.eta = eta  
        self.max_iter = max_iter  
        self.random_state = random_state  
        self.w = None  
        self.costs = [ ]
```

```
    def activation(self, X):  
        return self.net_input(X)
```

```
    def predict(self, X):  
        return np.where(self.activation(X) ≥ 0.0, 1, -1)
```


```

def fit(self, X, y):
    rgen = np.random.RandomState(self.random_state)
    self.w = rgen.normal(loc = 0.0, scale = 0.01, size = 1 + X.shape[1])
    self.costs = [ ]
    for n_iter in range (self.max_iter):
        idx = rgen.permutation(len(y))
        X, y = X[idx], y[idx]
        cost = 0
        for xi, yi in zip(X, y):
            error = yi - self.predict(xi)
            self.w[1:] += self.eta * error * xi
            self.w[0] += self.eta * error
            cost += error**2
        cost /= 2
    self.costs.append(cost)


```

# Table of contents

- 1 Perceptron
- 2 Linear Regression
- 3 **Adaptive Linear Neuron (Adaline)**
- 4 Logistic Regression



**Adaline**



**Logistic  
Regression**

**Figure:** Differences between Adaline and Logistic regression

# Components

## Hyperparameters

- eta
- max\_iter
- random\_state

## Parameters

- $w$
- costs

## Methods

- $\text{fit}(X, y)$
- $\text{predict}(X)$
- $\text{net\_input}(X)$
- $\text{activation}(X)$

# Implement (code from scratch)

```
class LogisticRegression:
```

```
    def __init__(self, eta = 0.01, max_iter = 50, random_state = 1):  
        self.eta = eta  
        self.max_iter = max_iter  
        self.random_state = random_state  
        self.w = None  
        self.costs = [ ]
```

```
    def activation(self, X):  
        return 1. / (1. + np.exp(-np.clip(self.net_input(X), -250, 250)))
```

```
    def predict(self, X):  
        return np.where(self.activation(X) ≥ 0.5, 1, 0)
```

```

def fit(self, X, y):
    rgen = np.random.RandomState(self.random_state)
    self.w = rgen.normal(loc = 0.0, scale = 0.01, size = 1 + X.shape[1])
    self.costs = [ ]
    for n_iter in range (self.max_iter):
        output = self.activation(X)
        errors = y - output
        self.w[1:] += self.eta * X.T.dot(errors)
        self.w[0] += self.eta * errors.sum()
        cost = (-y.dot(np.log(output)) - ((1 - y).dot(np.log(1 - output))))
        self.costs.append(cost)

```



# Practice

```
>> clf_LR = LogisticRegression(eta=0.01, max_iter=20,  
    random_state=1)  
    clf_LR.fit_SGD(X_train, y_train)  
  
>> y_pred = clf_LR.predict(X_test)
```

# Implement (library)

## Syntax (import)

```
from sklearn.linear_model import LogisticRegression
```

## Examples

```
>> from sklearn.linear_model import LogisticRegression as  
    LogisticRegression_  
>> clf_LR_lib = LogisticRegression_(random_state=1)  
    clf_LR_lib.fit(X_train, y_train)  
>> y_pred_lib1 = clf_LR_lib.predict(X_test)
```